

Ben Bubnick

HEALTHCARE DATA STRATEGY · DATA GOVERNANCE · ANALYTICS CONSULTING

✉ ben.bubnick@gmail.com | 📱 bubnicbf | 🌐 benbubnick
(216) 556-1612

Summary

Experienced Data Engineer with a strong track record of delivering scalable health data pipelines across cloud platforms, automating validation processes, and documenting workflows.

Experience

HHaEXchange

Snowflake, dbt, SQL/Jinja, Python, Snowpark, pandas, Liquibase/GitHub Actions, FHIR/Aidbox

DATA ENGINEERING CONSULTANT — CONTRACT VIA GROWDATA ANALYTICS

Oct. 2025 - Jun. 2026

Led data engineering work spanning source analysis, data modeling, ETL development, identity resolution, and QA to deliver standardized healthcare datasets for reporting and downstream integration.

- Led delivery of Snowflake/dbt pipelines that transformed operational healthcare data into standardized, analytics-ready models.
- Designed organization, office, caregiver, patient, and crosswalk models using FHIR-aligned concepts and stable identity-resolution patterns.
- Built reusable SQL/Jinja transformations for lifecycle fields, null handling, sentinel dates, and source-system variations.
- Developed QA suites covering duplicates, nulls, referential integrity, business rules, and source-to-target reconciliation.
- Partnered with stakeholders to resolve source ambiguity, document mapping standards, and establish reusable onboarding patterns.

Adonis

AWS: S3, Snowflake, SQLMesh, SQL, Python

DATA ENGINEERING CONSULTANT — CONTRACT VIA HEALTH DATA MOVERS

Mar. 2025 - Jul. 2025

Developed a reusable ETL pipeline to streamline data onboarding for a healthcare revenue cycle platform.

- Built a reusable SQLMesh/Snowflake ingestion framework with templated models, data contracts, CI checks, and promotion gates.
- Added schema-drift detection, reconciliation checks, audit logs, and deterministic recovery procedures.
- Developed modular transformations and validation checkpoints that reduced regression risk and accelerated onboarding.
- Translated client requirements into technical specifications, mapping standards, and reusable integration playbooks.

Arkos Health

AWS: S3, Athena, Redshift, Glue, dbt, PySpark, Docker + AWS CLI, Tableau, Jupyter

SENIOR DATA ANALYST

Jun. 2023 - Jan. 2025

Built a scalable AWS data lake with automated CI/CD, metadata reporting, and integrated analytics tools.

- Built an AWS data lake with dbt-based CI/CD, cost-aware partitioning, and job scheduling, cutting deployment time by approximately 50%.
- Improved reliability through pipeline monitoring, failure classification, alerting, and incident-response runbooks.
- Integrated Jupyter-based reporting and documented the platform, improving analytical usability and developer onboarding.
- Mentored engineers and established standards for code review, testing, pipeline design, and stakeholder delivery.

Merative

Azure: Data Factory, Blob Storage, Functions, SQL DB, Synapse Analytics, Python, DevOps

SENIOR DATA SCIENTIST

Jul. 2022 - Mar. 2023

Launched a new Azure platform to replace HDFS systems to modernize data infrastructure.

- Led HIPAA-compliant deidentification across more than 200 million patient records.
- Built automated regression and data-comparison tests to detect schema drift and transformation errors.
- Validated migration workflows across Azure Data Factory, Blob Storage, SQL Database, and Synapse Analytics.
- Documented migration procedures and enforced governance controls supporting HITRUST readiness.

IBM

Hadoop/HDFS/CDH: Spark, Hive & Impala, Kafka, Python, JRuby, SQL, Pig, Docker + Gradle

SENIOR DATA SCIENTIST

Nov. 2021 - Jul. 2022

Built AI tools while leading analytics delivery and large-scale data validation efforts.

- Built surgical scheduling optimization and CNN/RNN case-card solutions.
- Led client analytics delivery and validated the migration of more than 150 billion patient records.

SENIOR DATA ARCHITECT

Dec. 2020 - Nov. 2021

Led large-scale clinical data and NLP initiatives and delivering ML solutions.

- Led integration of more than 100 billion clinical records and topic modeling across nearly six million provider notes.
- Built a COVID-19 hospitalization model while documenting AI workflows and mentoring new hires.

DATA ARCHITECT

Nov. 2017 - Dec. 2020

Led enterprise data integration initiatives and mentoring programs.

- Managed multi-org enterprise integration with over 70 interconnected data sources.
- Developed new data validation and review process that reduced total rework by 25%.

SOFTWARE DEVELOPER, DATA PLATFORM

Sep. 2016 - Nov. 2017

Developed internal resources that reduced costs and improved onboarding efficiency.

- Built a data-curation tool that replaced a \$30,000 annual contract.
- Redesigned documentation and training, reducing knowledge-transfer time by 80%.

DATA SCIENTIST

Nov. 2015 - Sep. 2016

Led integration efforts and process automation initiatives.

- Managed 12 remote contractors across 36 data-integration projects.
- Built automation that eliminated 28% of unnecessary development work.

Explorys

Hadoop/HDFS: Ruby, Pig, React, NodeJS, JavaScript, SQL, Oracle, Jenkins

DATA SCIENCE CONSULTANT — CONTRACT VIA TEKSYSTEMS

Jun. 2015 - Nov. 2015

Built ETL pipelines on distributed data systems to support clinical data integration and population analytics.

- Served as principal integration specialist for more than 10 client data warehouses.
- Troubleshoot production pipelines and built scheduling controls for client data ingestion.

Other Experience

PRINCIPAL ML ENGINEER, **GrowData Analytics**, *Sep. 2024 - Oct. 2025.*

SOFTWARE ENGINEER, **AmTrust Financial Services**, *Mar. 2014 - Jun. 2015.*

COLOR SCIENTIST, **PPG Industries**, *Apr. 2012 - Jun. 2013.*

ADJUNCT PROFESSOR, **Multiple Institutions**, *Aug. 2010 - May. 2019.*

Education

2022 **PGCert.**, Data Science Professional

IBM

2015 **PGCert.**, Data Science

Johns Hopkins University

2010 **Master of Science**, Astrophysics

University of Cincinnati

2004 **Bachelor of Fine Arts**, Printmaking

Ohio University