

Ben Bubnick

DATA ENGINEER · DATA SCIENTIST · ANALYTICS SPECIALIST

✉ ben.bubnick@gmail.com | 📱 bubnicbf | 🌐 benbubnick
(216) 556-1612

Summary

Experienced Data Engineer with a strong track record of delivering scalable health data pipelines across cloud platforms, automating validation processes, and documenting workflows.

Experience

GrowData Analytics

AWS: S3, Glue, Athena, Lambda, CloudWatch, dbt, Snowflake, Ruby, RubyLLM, Tesseract OCR, GPT 4o mini, Cladue Opus/Sonnet, Rubocop, RSpec

PRINCIPAL DATA ENGINEER

Oct. 2025 - Present

Led data engineering work spanning source analysis, data modeling, ETL development, identity resolution, and QA to deliver standardized healthcare datasets for reporting and downstream integration.

- Led end-to-end design and delivery of healthcare data integration pipelines, converting operational source data into curated analytics-ready models for reporting and downstream use.
- Designed and refined core domain models for organization, office, caregiver, patient, and crosswalk entities to support standardized healthcare data workflows.
- Defined surrogate key, natural key, and crosswalk strategies to enable stable identity resolution across related source systems.
- Designed and executed comprehensive QA and data-integrity test packs covering row counts, duplicates, nulls, referential integrity, reconciliation, and exception handling.
- Collaborated with client and internal stakeholders to resolve source ambiguity, refine requirements, and translate business needs into production-ready data assets.

PRINCIPAL ML ENGINEER

Sep. 2024 - Oct. 2025

Built and shipped an OCR-to-analytics product for small farms; owned model integration, CI/CD, testing, and auditability on UNIX-like systems.

- Productized OCR → structured data pipeline for farm invoices (handwritten/printed) using Tesseract + rule-based post-processing; shipped as a small Python/Ruby service with containerized deployments and versioned configs.
- Built validation & reconciliation layers (schema checks, cross-field rules, idempotent re-runs) and produced audit logs to support bookkeeping/compliance workflows.
- Implemented CI/CD (unit tests, linting, fixtures, golden-file tests) and release automation; documented runbooks and failure modes (timeouts, low-confidence OCR, retries).
- Delivered a lightweight analytics layer (usage metrics, error taxonomy, latency/throughput dashboards) to guide feature tuning and cost control.
- Owned infra as code for a minimal AWS footprint (S3 for artifacts, Lambda/Cron for jobs, CloudWatch alerts); operated via UNIX-like CLI workflows (Bash/Zsh, Makefile targets).

Health Data Movers

AWS: S3, Snowflake, SQLMesh, Python, SQL, Docker

DATA CONSULTANT

Mar. 2025 - Jul. 2025

Developed a reusable ETL pipeline to streamline data onboarding for a healthcare revenue cycle platform.

- Built a reusable ingestion & transformation framework in SQLMesh/Snowflake (templated models, data contracts, CI checks, promotion gates), enabling consistent, repeatable onboardings.
- Added observability & auditability (row-count/parity checks, schema drift alerts, reconciliation logs) and wrote runbooks for deterministic re-runs and incident handling on UNIX-like systems.
- Shipped a lightweight forecasting/alerting pattern (Python service + Docker + Actions) for operations signals; packaged with tests and deployment automation.
- Created testable, modular SQLMesh models, reducing regression errors and accelerating iteration cycles.

Arkos Health

AWS: S3, Athena, Redshift, Glue, dbt, PySpark, Docker + AWS CLI, Tableau, Jupyter

SENIOR DATA ANALYST

Jun. 2023 - Jan. 2025

Built a scalable AWS data lake with automated CI/CD, metadata reporting, and integrated analytics tools.

- Stood up an AWS data lake with CI/CD (dbt + Git flows), cost-aware partitioning, and job scheduling; reduced deployment time 50% and improved query performance with targeted model refactors.
- Implemented monitoring/alerting & on-call runbooks (metrics, logs, failure taxonomies) to lower MTTR and improve data reliability.
- Mentored engineers on code reviews, testing, and pipeline design; authored internal standards that lifted team delivery quality.
- Integrated Jupyter into reporting, boosting report interactivity and increasing stakeholder satisfaction by 30%.

Merative

Azure: Data Factory, Blob Storage, Functions, SQL DB, Synapse Analytics, Python, DevOps

SENIOR DATA SCIENTIST

Jul. 2022 - Mar. 2023

Launched a new Azure platform to replace HDFS systems to modernize data infrastructure.

- Led HIPAA-compliant deidentification initiative across 200+ million patient records.
- Designed and implemented an automated testing framework for ETL pipelines for validation coverage.
- Validated and optimized Azure-based data workflows, ensuring continuity and fidelity during the transition from legacy HDFS systems to cloud architecture.
- Conducted end-to-end system testing for Azure Data Factory pipelines, Blob Storage configurations, and downstream Synapse integrations.

IBM

Python, GoLang, Azure, Docker, SQL/Hive/Impala, Tableau, Ruby, Pig, kafka, Kibana,

Hadoop/HDFS, Gradle, Spark, Perl

SENIOR DATA SCIENTIST

Jun. 2015 - Jul. 2022

Built AI tools while leading analytics delivery and large-scale data validation efforts.

- Created AI Surgical Scheduling Optimization tool currently evolving into new offering.
- Developed clinical data integration of 100+ billion structured & semi-structured records.
- Developed surgical smart case card using combined CNN/RNN model.
- Developed topic modelling project on nearly 6 million semi-structured provider notes data.

Explorys

Hadoop/HDFS: Ruby, Pig, React, NodeJS, JavaScript, SQL, Oracle, Jenkins

DATA SCIENTIST

Jun. 2015 - Nov. 2015

Built ETL pipelines on distributed data systems to support clinical data integration and population analytics.

- Principal integration specialist for 10+ client database warehouses.
- Performed troubleshooting on production issues across multiple environments.

Other Experience

ADJUNCT PROFESSOR, **Multiple Institutions**, Aug. 2010 - May. 2019.

SOFTWARE ENGINEER, **AmTrust Financial Services**, Mar. 2014 - Jun. 2015.

COLOR SCIENTIST, **PPG Industries**, Apr. 2012 - Jun. 2013.

Education

2022 **PGCert.**, Data Science Professional

IBM

2015 **PGCert.**, Data Science

Johns Hopkins University

2010 **Master of Science**, Astrophysics

University of Cincinnati

2004 **Bachelor of Fine Arts**, Printmaking

Ohio University